

Charlie Crisp

Building a Blockchain Library for OCaml

Computer Science Tripos – Part II

Pembroke College

March 26, 2018

Proforma

Name: **Charlie Crisp**
College: **Pembroke College**
Project Title: **Building a Blockchain Library for OCaml**
Examination: **Computer Science Tripos – Part II, July 2018**
Word Count: **????¹**
Project Originator: **KC Sivaramakrishnan**
Supervisor: **KC Sivaramakrishnan**

Original Aims of the Project

To build a library in OCaml, which can be used as a building block for blockchain applications. The library should allow participating nodes to own a shared copy of a blockchain data structure, agreed upon using consensus. Nodes should also be able to commit transactions to the blockchain, which should then be visible to other participating nodes.

Work Completed

All that has been completed appears in this dissertation.

Special Difficulties

None

Declaration

I, Charlie Crisp of Pembroke College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Signed

Date

Contents

1	Introduction	9
1.1	The History of the Blockchain	9
1.2	Blockchain Today	10
1.3	Work Completed	10
2	Preparation	11
2.1	Starting Point	11
2.2	Using OCaml	11
2.3	Requirements Analysis	15
2.3.1	Data Structure	15
2.3.2	Consensus	17
3	Implementation	19
3.1	System Architecture	19
3.2	The Blockchain Data Structure	20
3.2.1	Irmin	20
3.2.2	Ezirmin	22
3.3	Consensus Algorithms	25
3.3.1	Building Consensus	26
3.3.2	A New Approach	28
4	Evaluation	35
4.1	Performance on a Local Machine	35
4.1.1	Single Participant	35
4.1.2	Multiple Participants	37
4.2	Performance on Remote Machines	37
4.2.1	Single Remote Participant	37
4.2.2	Multiple Remote Participants	37
5	Conclusion	39
	Bibliography	41
A	Project Proposal	43

List of Figures

1.1	A typical blockchain structure	10
3.1	Logan system architecture overview	20
3.2	Architecture of an Irmin Store	21
3.3	Merging Ezirmin Logs	24
3.4	Typical Participant mempool	31
3.5	Naïve approach to retrieving mempool updates	32
3.6	Missing mempool updates with a naïve consensus mechanism	32
3.7	Advanced approach to retrieving mempool updates	34
4.1	Latency variation with throughput on a single local Participant	36
4.2	Latency variation with blockchain size on a local Participant	37
4.3	Logan performance with a single remote Participant	38

Acknowledgements

I would like to thank KC Sivaramakrishnan for being an extremely helpful supervisor throughout the duration of the dissertation, as well as over the past three years.

I would also like to thank Anil Madhavapeddy for allowing me to use his laptop for the duration of the dissertation, and being a very supportive DoS.

Finally I would like to thank my friends and family for supporting me through my final year.

Chapter 1

Introduction

Blockchain technology has existed for a long time, but the definition of ‘blockchain’ has changed drastically since its conception. Previously used just to describe a data structure, the term ‘blockchain’ is now widely used to also describe the accompanying consensus mechanisms. This is mainly due to the increasing popularity of cryptocurrencies such as Bitcoin [9] which use the ‘Proof of Work’ algorithm to solve the double spending problem. Blockchain is undoubtedly the most important technology in the field of cryptocurrencies, where no single client can be trusted, however, it also has many use cases outside this sector. For instance, it can be used in situations where clients *can* be trusted like a hospital maintaining internal medical records, or a bank wishing to record transactions from many of its own distributed clients.

In this report I will present ‘Logan’, a blockchain library in OCaml which allows the easy creation of blockchain applications. The blockchain is synchronised via a leader-based consensus mechanism with strict consistency. Developers using this library are also able to define custom validation of transactions being added to the blockchain. Because the application is written in OCaml, it can be compiled to bytecode, unikernels or even javascript and is therefore suitable for a wide range of destination applications and devices.

1.1 The History of the Blockchain

The blockchain, in its simplest form, is a series of blocks of data, where each block contains the hash of the content of the previous block. Figure 1.1 is a graphical representation of a typical blockchain data structure.

The blockchain, as a cryptographically secure chain of blocks, was first conceptualised by Stuart Haber and W. Scott Stornetta in 1990 [7]. However, until the creation of Git [12] in 2005, the blockchain was still a relatively niche concept. The invention of Bitcoin in 2008 is seen by many as the most pivotal moment in the history of blockchain technologies. Bitcoin uses the Proof of Work consensus algorithm to create a decentralised, trust-less, peer to peer network which is used to make transactions between virtual wallets.

Chapter 2

Preparation

Before starting work on the project code, I completed a lot of preparation in order to aid the development process. I spent some time learning to use OCaml and familiarising myself with its features. This was important because it allowed me to write idiomatic code which was not only powerful, but could also provide an intuitive API for other developers to use. I also spent some time investigating a few key libraries, such as Irmin and Lwt. Understanding these libraries, the data structures they provide, and the technologies they present, allowed me to focus on the technological challenges in the project and not waste time reinventing the wheel. Setting up a good development environment allowed me to run automated builds and, therefore, catch any errors in the code early. This was particularly useful during the evaluation as it involved installing the project on many remote machines. Being able to catch issues like dependency resolution or out of date build commands made the evaluation process much easier. Lastly, I spent some time developing a requirements analysis for the final product. This analysis helped drive the design and development of the blockchain library whilst not restricting the work that I was able to complete.

2.1 Starting Point

The project built upon functionality provided by Irmin [1] which is a distributed database system. Irmin is fast, durable and has all the necessary capabilities required to build a blockchain. The project also made use of Ezirmin [11] which provides a simplified API to Irmin, and defines a log data structure which was used to build the blockchain and mempool structures for this project.

2.2 Using OCaml

At the start of the project, I had never used OCaml for any project of significance. Whilst the first year Foundations of Computer Science course had given me some background into functional programming, there were still many key OCaml features which I had to learn. In the first few weeks of my project, I spent time studying the book Real World OCaml [13] which proved a great introduction to many of OCaml's features.

Pattern Matching

OCaml provides a very powerful syntax for matching patterns which allow you to write functions as shown by Listing 2.1.

Listing 2.1: OCaml Pattern Matching

```
let pattern_matcher_1 input = match input with
| Card("spades", 1) -> Printf.printf "It's the ace of
    spades!"
| _ -> Printf.printf "Unlucky"
let pattern_matcher_2 = function
| Card(suit, 1) -> Printf.printf "It's the ace of %s!"
    suit
| _ -> Printf.printf "Unlucky"
```

Listing 2.1 shows a function that will print a special string if it is passed the Ace of Spades. Here, the pattern matching checks that the tuple associated with the data type contains the string `spades` and the number `1`. The second example assigns the variable `pattern_matcher_2` to an anonymous function which matches the first argument of the tuple to the variable `suit`.

Options

An `Option` is a built in data type that provides a powerful way of using OCaml's pattern matching. By using the `Some (...)` and `None` constructors, one can create something of the type `'a option`. This is comparable to the `null` type in languages such as Java, however as it is part of the type system, it forces the programmer to handle any cases where `null` could be returned.

Error Handling

OCaml provides multiple different ways of dealing with errors and exceptions. A simple way of signifying an error in your return type, is to return an `option` which will be `None` if there is an error. Whilst this can be inflexible for larger solutions, it also provides a quick and simple way of signifying that something has gone wrong.

Another way of dealing with options in return types is to use the `bind` function:

```
val bind: 'a option -> ('a -> 'b option) -> 'b option = <fun
    >
```

As the type signature demonstrates, `bind` takes an `option` and applies a function to its contents if it exists, or returns `None` otherwise.

OCaml provides a built in type `Result.t` which is an extension of optional return types, where the programmer is able to define arbitrary data to accompany the error type. The following demonstrates a successful return type of `int` (with an unspecified `Error` type), and a `string` error type (with an unspecified `Ok` type).

```
# Ok 3;;
- : ('a, int) result = Ok 3
# Error "Something went wrong";;
- : (string, 'a) result = Error "Something went wrong";;
```

Polymorphic Variants

OCaml allows the programmer to define variant types such as the `Card` type that was defined earlier. This makes it very easy to make use of pattern matching with custom defined types.

OCaml also introduces the notion of polymorphic variants which are more flexible and do not require an explicit type declaration.

Listing 2.2: Polymorphic Variants

```
# let card = `Card ("spades", 1);;
- : val card : [> `Card of string * int ] = `Card ("spades",
  1)
```

Listing 2.2 shows how a back-tick can be used to define a polymorphic type, and OCaml has automatically inferred a type. The `>` symbol acts as a lower limit on the tags that the variant `card` can take, i.e. it can have the tag `Card` or any other unspecified tag. When dealing with variant types as *parameters*, the `<` symbol often appears in the type signature to show that the parameter can belong to only a subset of given set of tags. The absence of both of these symbols indicates that a variant can take any value from the given set of tags, rather than a superset or subset.

Modules and Functors

Modules provide a way of grouping together related code in OCaml. They can be thought of as similar to traditional namespaces, although there are a few key differences. OCaml also lets you define module type signatures which modules have to conform to.

Listing 2.3: OCaml Modules and Functors

```
module type Math = sig
  type t
  val add: t -> t -> t
  val subtract: t -> t -> t
end
```

```

module IntegerMath : Math with type t = int = struct
  type t = int
  let add int1 int2 = int1 + int2
  let subtract int1 int2 = int1 - int2
end

```

Listing 2.3 is a simple example that defines a module signature `Math` for adding and subtracting a custom type. The module `IntegerMath` is a module which adheres to this signature. The syntax `with type t = int` serves the purpose of letting the compiler know that the type `t` is externally visible. Whilst this example uses the trivial example of integer maths, it is easy to imagine this extending to, for example, matrices or sets where these functions are not part of the language.

Modules are really useful for allowing the effective division of code into isolated units, but on their own, they are slightly inflexible. Maybe a developer would want to abstract some lower level details of code used by a module. In this case, she would have to create a whole new module for each possible implementation of this abstraction. An example of this could be a datastore that uses either an in-memory or on-disc format for storing data. Functors allow modules to be created from other modules.

Listing 2.4: Ezirmin Log Module

```

module Log(AO : Irmin.AO_MAKER)(V: Tc.S0) = struct
  ...
end

```

Listing 2.4 is an example from the Ezirmin codebase and shows how to define a functor. The parameter `AO` is a module used for creating append only stores, and the parameter module `V` defines a data type. The result of this is a functor, `Log`, which can be used to create a `Log` module with either an in-memory or on-disc backend.

Development Environment

When developing a large scale system with OCaml, there are a couple of build systems available to use. ‘jbuilder’ [4] is one of these systems which is becoming increasingly popular and is used daily by hundreds of developers. jbuilder allows the developer to specify arbitrary directory structures containing executables, libraries and more. I set up my project to compile Logan, which includes the interface for running both Leader and Participant nodes. I also defined multiple executables for running both the Leader and Participants in example and test cases. I used GNU Make [3] to invoke jbuilder which allowed me to easily build and run any executables from the root of the directory.

In order to ensure that the project would always build, I set up a continuous integration workflow using Travis-CI. This ensured that whenever I pushed any updates to my GitHub

repository, Travis would attempt to build the system and would notify me whenever there were any errors during the build. This was particularly useful during the Evaluation stage when I was installing my project on multiple remote machines for testing. Over the development period, Travis flagged issues such as dependency resolution failures or out of date build scripts and I was able to fix these at the point when they first occurred, rather than weeks later when they would have caused many issues during remote installation.

Lwt

A library which I spent some time familiarising myself with was Lwt [5]. Lwt provides a way of interacting with threads in OCaml, although in Lwt they are known as ‘Promises’.

Listing 2.5: Lwt Promises

```
val Lwt.return : 'a -> 'a Lwt.t
val Lwt_main.run : 'a Lwt.t -> 'a
val Lwt.bind : 'a Lwt.t -> ('a -> 'b Lwt.t) -> 'b Lwt.t
```

Listing 2.5 shows the basic functions for creating, running and combining threads. The above type `'a Lwt.t` refers to a thread which will eventually terminate with a value of type `'a`, and follows the well established Monad design pattern. `Lwt.return` is a function that takes a value and will create a promise that immediately returns with this value. This is useful when inserting static or precomputed variables into Lwt pipelines. `Lwt_main.run` is the dual of `Lst.return` and will run a Lwt promise until completion and then return its value. This is useful when retrieving values from an Lwt pipeline. Finally, `Lwt.bind` (or the infix notation `>>=`) will pass the result of a Lwt promise to a function returning another Lwt promise. This is useful for chaining together Lwt promises in a pipeline.

2.3 Requirements Analysis

During the preparation stage of my project, I spent some time analysing the requirements that would be suitable for Logan. This proved a good way of guiding the progress of the project and making sure that I solved all the problems that I set out to. Here, I will set out the criteria that I decided upon before starting development on the project.

2.3.1 Data Structure

One of the most important requirements for Logan, is that it should be able to add items to and view items in a blockchain. The data in the blockchain will be referred to as either `blocks` or `transactions`. From a participating node, it should be possible to add a transaction, of arbitrary type. It should also be possible to view an ordered list of all transactions which currently exist in the blockchain. The latency of adding transactions should not increase as the size of the blockchain increases because it is

important the library will cope well with systems running for an indefinite period of time.

Listing 2.6: Blockchain Specification

```

module type I_LogStringCoder = sig
  type t
  val encode_string: t -> string
  val decode_string: string -> t option
end

module type I_Config = sig
  type t
  module LogCoder: I_LogStringCoder with type t = t

  val validator: (t list -> t -> bool) option
end

module type I_Blockchain = sig
  type t

  val add_transaction_to_blockchain: t -> [> `Error | `Ok]
    Lwt.t
  val get_all_transactions: unit -> [> `Error | `Ok of t
    list] Lwt.t
  val get_transactions: int -> [> `Error | `Ok of t list]
    Lwt.t
end

module Make(Config: I_Config): I_Blockchain with type t =
  Config.t = struct
    ...
  end

```

Listing 2.6 is the initial technical specification that I used to define Logan, and the functions complete the following operations:

- `I_LogStringCoder` is a module that allows the user to specify arbitrary types to be stored on the blockchain, so long as they can be encoded to (and decoded from) a string.
- `I_Config` contains information which is required to run the blockchain. In particular, should the `validator` option contain the value `Some(f)`, then `f` will be a function that accepts a history of committed transactions, and validates whether a further transaction is valid.
- `Make` is a functor which accepts a configuration module and will return a blockchain module.

- `add_transaction_to_blockchain` adds a user defined type to the blockchain and then returns a polymorphic variant type containing information about whether the operation was successful. This will return an error in the case that the transaction is not validated.

This specification differs slightly from the final specification for Logan because it does not take into account the difference between Leader and Participant nodes. The specification was deliberately designed without focus on consensus so that the focus was on the external functionality rather than the internal implementation. The final interface, while different, provides exactly this functionality.

2.3.2 Consensus

Ensuring that Logan maintains consensus was a very crucial part of this project. The actual design and development of the consensus algorithm was completed throughout the duration of the project and involved a lot of research into other consensus mechanisms. However, during the requirements analysis phase of the project I set out some goals for the final implementation. These goals were laid out in order to help drive the design and development of the consensus mechanism. I decided on the following requirements:

- The consensus mechanism must guarantee strict consistency. Although strict consistency could add large overhead costs, it is necessary for some applications and so is an important part of the final project.
- The consensus mechanism must be scalable. Within the scope of this project, it should be possible for the blockchain to be shared by 3 or more nodes in a network. This should not hinder the performance of the system, and it should still be able to handle multiple successful transactions per second.

Chapter 3

Implementation

Over the course of this project, I have successfully built Logan, a blockchain library which can perform custom transaction validation. Logan uses a leader based consensus protocol inspired by Raft also uses the concept of a ‘transaction waiting room’, or *mempool*, which is used by Bitcoin. Section 3.1 gives a general overview of Logan’s system architecture. Section 3.2 describes the blockchain data structure that I have used, giving a justification of why it can be considered a blockchain. This section also gives a description of the merge semantics of Ezirmin Logs, which are important to understand when describing the behaviour of the Leader mempool. I will also describe two bugs that I encountered in these codebases and the fixes that I implemented for each. Finally, Section 3.3 gives a description of the research I completed when designing Logan’s consensus mechanism. I will first present a naive approach to consensus and explain why this is flawed and will actually ‘miss’ transactions. I will then present an improved mechanism, which Logan uses, that guarantees that no transactions in participant mempools are missed.

3.1 System Architecture

Logan goal is to allow multiple nodes in a distributed system to share a blockchain data structure. Nodes in this system should be able to add transactions to the blockchain and read the entire history of previous transactions. Logan aims to achieve this goal by using a Leader-based approach where a ‘Leader’ node can accept transactions from ‘Participant’ or ‘Worker’ nodes and commit them to the blockchain for everyone to see. Participants view the blockchain by retrieving the latest copy over an SSH connection with the Leader. Participants request for transactions to be committed by writing to their local mempool. The Leader periodically polls the Participant mempools, again over an SSH connection, to find new transactions to add to the blockchain. When a poll finds new transactions, it will either commit them all, or in the case that a validator function is supplied, it will retrieve all current items in the blockchain and use these to filter out invalid transactions before committing. It is also possible for a Participant node to exist on the same physical machine as the Leader, in which case the leader sources updates directly from the mempool, rather than over an ssh connection. Figure 3.1 shows how the architecture for this system is organised.

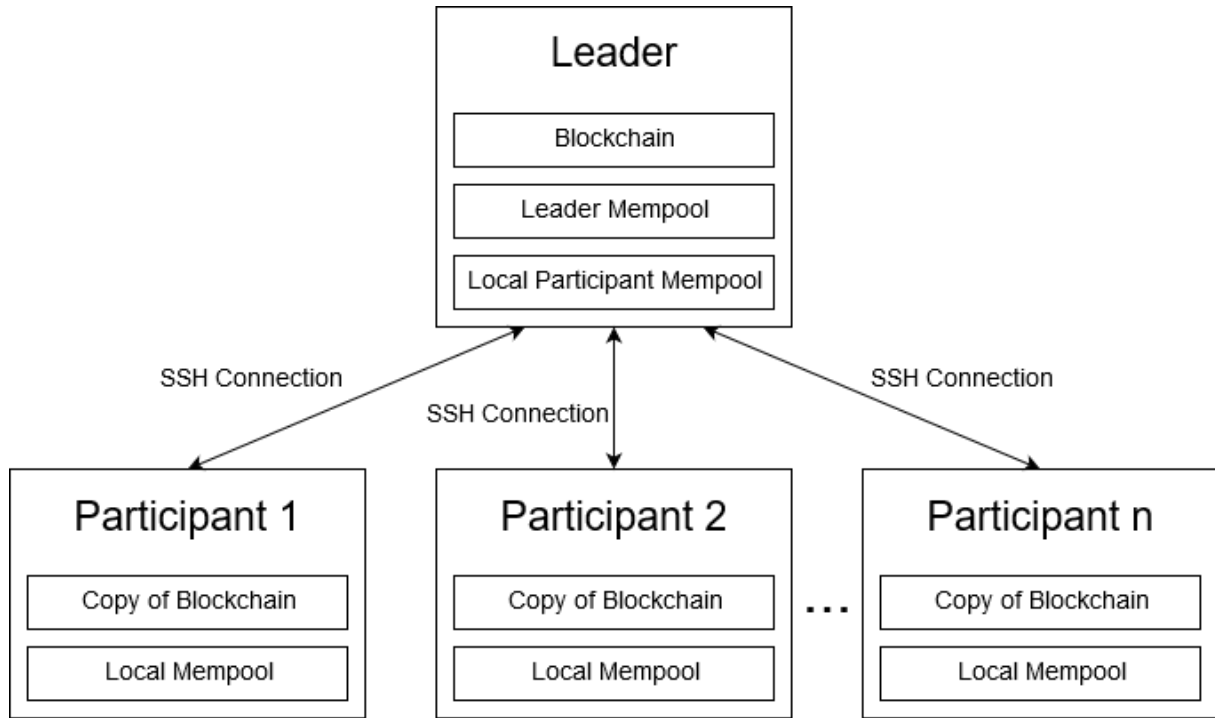


Figure 3.1: Logan system architecture overview

3.2 The Blockchain Data Structure

Irmin is a library for OCaml providing datastore functionality using a Git backend. Ezirmin is a wrapper around Irmin that provides a simple log data structure which I have used as a blockchain. In order to justify that this data structure can be considered a blockchain, I took some time to investigate its semantic properties. Whilst there is no universally agreed-upon definition of a blockchain, I have used the following criteria to define the blockchain:

1. Data is stored in ‘blocks’.
2. Blocks are ordered in a tree structure where each block contains the hash its parent block.

This definition deliberately makes no mention of consensus, and exists purely to justify the validity of the blockchain data structure.

3.2.1 Irmin

Irmin is a library for OCaml which provides Git-like, distributed, branchable storage [6]. Irmin exposes three main structures as shown in Figure 3.2: The Block Store, the Tag Store and the Irmin Store. For the reader familiar with Git, these can intuitively be thought of as objects/commits, branches and repositories respectively.

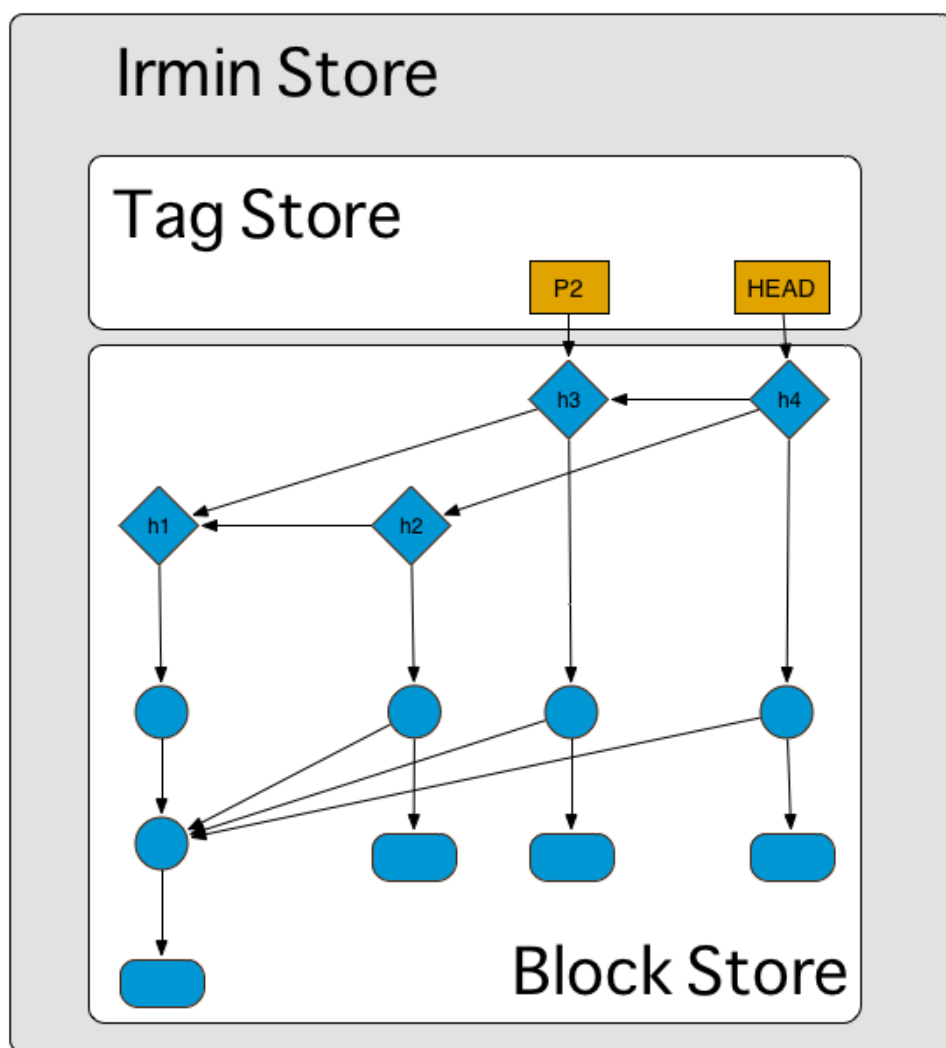


Figure 3.2: Architecture of an Irmin Store

The Block Store

The Irmin Block Store is a heap of immutable blocks. Rather than being addressed by a physical address, these blocks are addressed by the hash of their content. Because the Block Store is content-addressable, once blocks are added, their content can never be updated. Instead, updates to Irmin data structures will add new blocks which try to utilise as much shared history with the existing data in the store, in order to minimise storage space.

The Tag Store

Any immutability in Irmin data structures derive from *tags*. Tags provide a way of indexing into any block in the Block Store. This concept is similar to that of branches or references in Git, which provide a way of indexing into a particular commit in the Git history. Because blocks are immutable, any changes to an Irmin data structure can only be visible if a tag is updated. In particular, I will refer to the HEAD tag or pointer, which will index into the latest recognised log item. When Irmin merges data structures, it creates new blocks and updates the relevant tags to point to these, however, the old blocks remain in the block store.

Irmin Stores

An Irmin Store is simply the combination of a Block Store and a Tag Store. Considering the criteria set out earlier for a data structure to be considered a blockchain, Irmin Stores satisfy the first criteria, i.e. that data is stored in Blocks. Indeed, they also come close to satisfying the second criteria, as all blocks are indexed by the hash of their content. However, this is not quite enough, as blocks are not required to contain hashes of parents in a chain. In order to satisfy these criteria, I looked to Ezirmin which provides a log data structure.

3.2.2 Ezirmin

Ezirmin is a library that provides a simplified interface to the Irmin library. It is designed to provide a interface to Irmin without functors, but with some useful defaults. Importantly, it has a built in log data structure which uses Irmin's append-only store, saved on disk in the Git format.

Listing 3.1: Ezirmin Log

```
module type FS_Log = sig
  type elt
  type cursor
  val append : ?message:string -> branch -> path:string list
    -> elt -> unit Lwt.t
  val get_cursor : branch -> path:string list -> cursor
```

```

    option Lwt.t
  val read : cursor -> num_items:int -> (elt list * cursor
    option) Lwt.t
  val read_all : branch -> path:string list -> elt list Lwt.
    t
  ...
end

```

Listing 3.1 gives some of the interface for an Ezirmin log which uses a file system backend. The log allows for items to be appended to and read from a log, and the concept of cursors here maps directly to the concept of Irmin tags. In particular, the function `read` will read from the item indexed by a cursor, and will return a new cursor for the next unread log item, alongside the result of the read.

Ezirmin Log as a Blockchain

Logan uses an Ezirmin Log as an implementation of a blockchain, but to verify that this is a valid data structure to use, I investigated the implementation of Ezirmin Logs. Ezirmin Logs use Irmin blocks to store log entries, and therefore satisfy the first criteria for being considered a blockchain. To verify that they also satisfy the second criteria, i.e. that blocks are ordered and contain the hash of their parent, I looked at the implementation of log items.

Listing 3.2: Ezirmin Log Item

```

type log_item =
{ time      : Time.t;
  message   : V.t;
  prev      : K.t option}

```

Listing 3.2 is taken from the Ezirmin Log implementation and shows how each log item has a `prev` value which is a key pointing to its parent log item. This imposes an ordering of log items, and means that any changes to previous log items will not be observed, as they will simply create a new block with a new address. Merging Ezirmin Logs causes new Merge blocks (as opposed to Value blocks) to be created that collate multiple log entries from divergent log histories. These entries are ordered by their timestamps and this means that after a merge, items may appear to have been inserted into the log history. This is not desirable behaviour for a blockchain, so I have ensured that the blockchain data structure maintained by Logan is never merged with other structures, only appended to. Consequently, each blockchain transaction will exist in its own Value block and will never removed or reordered in the blockchain. Therefore it can be concluded that an Ezirmin Log satisfies both of my conditions to be considered a blockchain.

Merging Ezirmin Logs

Aside from using an Ezirmin Log as the blockchain data structure for this project, I also use a *mempool* log to retrieve updates/transactions from remote mempools.

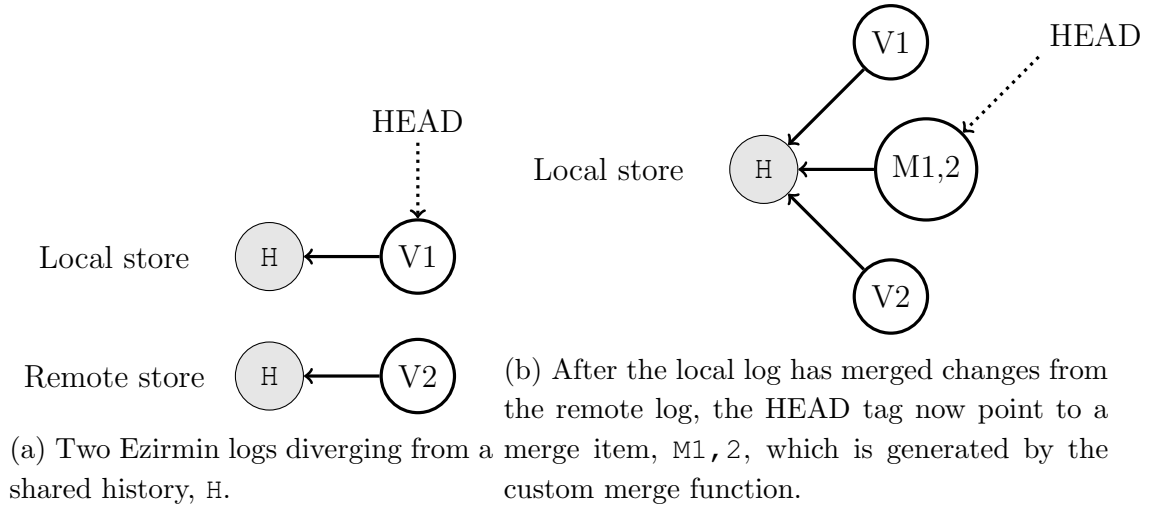


Figure 3.3: Merging Ezirmin Logs

This retrieval process uses the `Ezirmin.Repo.Sync` module to merge new log items into the Leader's mempool. Under the hood, this module uses the git protocol to pull updates over an SSH connection. In this section I will investigate the semantics of Ezirmin merges as they are critical to understanding how Logan achieves consensus.

Irmin allows for the developer to specify custom merge strategies. Whenever a Irmin Store merges changes from a remote Store, Irmin will do one of two things:

1. If the newly discovered blocks do not have a divergent history to the current store, they will simply be added to the Block Store. In other words, if only new blocks have been added in the remote store, these will be added to the local store.
2. If the remote and local stores have divergent histories, then a custom three way merge is made using the heads of each store and their latest common ancestor. This merge is defined by the developer.

Figure 3.3 shows how a local Ezirmin Log performs a merge from a remote log with a diverging history. Before the merge, the local store has a Value block, V1, but the remote store has a different Value block, V2. After the local store has merged changes from the remote store, a new Merge block, M1,2, has been created which contains both histories as defined by the custom merge function. Blocks V1 and V2 still exist in the block store, but the HEAD tag has been updated to point to M1,2.

Ezirmin Bugs

Whilst Ezirmin provides a set of very desirable semantic properties, it is not a widely used library. Consequently, during the course of the project, I encountered bugs in the Ezirmin codebase which I had to debug and fix. In this section I shall outline the two bugs that I encountered the changes I made to the Ezirmin codebase to fix them.

The first bug occurs when updates are merged into a log over a network.¹ Irmin uses Git as a backend and in order to merge updates over a network, the `git pull` command is used. In Git, objects are stored as blobs on the leaves of a tree and `git pull` works by retrieving all blobs at the leaf nodes of the tree for a given branch. However in Ezirmin, blobs can contain pointers to other blobs, and `git pull` will not know to retrieve these blob too. This will cause an error when Ezirmin tries to read these objects which do not actually exist locally. The solution that I implemented to this problem is to use an internal branch which uses an Irmin Store to explicitly track every blob in the log. Items are added to this Irmin Store on the internal branch every time they are added to the log. This operation inserts these blobs into the git tree such that if `git pull` is performed on the internal branch, all the blocks required by the master branch will be retrieved.

The second bug that I encountered derived from the merge behaviour of Ezirmin logs when there are more than two remote machines merging updates from each others' log.² The issue is that when a Merge block is created, it may contain log items which, on other machines, are stored as Value blocks or as different Merge blocks. The result is that Irmin sees these blocks as different, and Ezirmin does not perform any checks to detect or remove duplicate log items occurring in different contexts. This behaviour resulted from complex sequences of merge operations and so was difficult to reproduce, but in the worst case caused the size of the blockchain to grow exponentially with the number of merge operations performed. In one case, a log with 40 unique log items grew to be greater than 120,000 log items in size due to duplicates. I implemented a fix by removing duplicate log items when a merge operation takes place. Duplicates are found by comparing the log item timestamps, and whilst this does not guarantee their individuality, the timestamps have such significant precision to make this highly unlikely.

3.3 Consensus Algorithms

Building consensus was by far the most important part of work completed for this project. In order to guide the design of the consensus mechanism, I completed an extensive amount of research on existing algorithms. This section will briefly summarise this research and my conclusions about their suitability for Logan. The resulting design for consensus is a leader based approach, where Participants can commit transactions to a mempool. The Leader polls these mempools for updates, which are validated and then committed to the blockchain. This blockchain can be read by all Participants, and provides a definitive source of ordered, committed transactions.

¹The issue on GitHub can be found at <https://github.com/kayceesrk/ezirmin/pull/7>

²The issue on GitHub can be found at <https://github.com/kayceesrk/ezirmin/pull/8>

3.3.1 Building Consensus

Proof of Work

Proof of Work (PoW) is a deceptively simple consensus mechanism, used by most cryptocurrencies to avoid the double spending problem. Transactions are contained within blocks which can be broadcast out to the network of participating nodes. Whenever a block is received by a participating node, the node checks that the block contains a proof of computational work done. This proof usually takes the form of a random sequence of data (this is known as a nonce) appended to the end of the block, causing the block's hash to be prefixed with a set number of 0s. This acts as a proof of computational work, because the data appended to the end of a block can only be found by a brute force method called mining, but can also be verified easily by simply computing the block's hash.

Why is this useful? Well, this makes a guarantee on the validity of the blockchain based on the simple assumption that more than 50% of the workforce is genuine. That is, if we assume that the longest chain of blocks is the correct one, then in order to create a sequence of biased transactions, we would have to create a chain longer than the correct one. This would require an equal number of 'Proof of Work's, which, due to the random nature of block mining, would require more than 50% of the workforce. Whilst it may be possible to maintain an equally sized chain with less than 50% of the workforce for a short period of time, the chances of this decrease rapidly as time passes. All in all this means that the longer a block has been in the chain, the more likely it is that the block is valid.

Whilst this forms a very effective mechanism for achieving consensus, there are also some considerable downsides to using a PoW approach to consensus. Firstly, there is a huge amount of computational work wasted in the process of mining. The effect of this is energy consumption [1] and wastage to a level which can cause serious environmental harm. PoW also does not generalise well outside of the scope of cryptocurrencies. It assumes no trust in any participants which may not be a suitable model for an application. Additionally, it also assumes that miners can be rewarded, usually with cryptocurrency, but this incentive is ad hoc and may not exist in other applications.

Mempools

Mempools are an important part of the design of Bitcoin, and whilst they are not inherently linked to the Proof of Work consensus algorithm, they are worth investigating. When a Bitcoin transaction is made, it is first written into what is known as a Mempool. This transaction can then be seen by participating miners, who can then choose to put this in the next block that they mine. This is significant, as it provides a 'waiting room' for any transactions that have not yet been validated.

Proof of Stake

The Proof of Stake (PoS) algorithm is used by some cryptocurrencies and works by randomly allowing participants to create (or ‘forge’) a single block. However, the probability that a participant is chosen to ‘forge’ a block, is weighted by its stake, such that participants with higher stakes in the blockchain are more likely to be chosen to forge a block.

So, why is PoS desirable? By far the most convincing reason for using PoS over PoW, is that there is no need to waste lots of energy in the process of mining. This hugely reduces the environmental impacts of scaling a PoS network. Using PoS also allows trust to be distributed according to an arbitrary heuristic which can be desirable property in some applications.

One of the flaws of PoS is that it does not have such a strong deterrent against attacks. With PoW, attacks require huge amounts of computational power and it is likely that to create an attack, you would have to spend more on hardware than you would gain. PoS does not have this same built in mechanism, and so there have been many suggested schemes for increasing the safety of PoS networks. For example, it is possible that participants should need to pay some form of deposit before forging blocks, which can be slashed if they break any rules. PoS also suffers from the same problem as PoW in that it does not generalise well. It is another example of a consensus mechanism designed for networks with a strong notion of Stake and with minimal trust in any individual participant.

Paxos

Paxos is a family of consensus protocols which can be used to guarantee consistency in distributed systems. It was first proposed in a paper by Leslie Lamport in 1998 [8], although the paper was first submitted in 1990. Named after a fictitious civilisation living on the island of Paxos, the algorithm puts forward a way for any number of nodes to propose and agree on a value. Participating nodes belong to various roles, one of which is known as a ‘Proposer’ or Leader.

The main part of the algorithm is split into two sections, propose and accept. In the first stage, a Proposer decides that it wants to propose a value and then broadcasts out a ‘Prepare’ message to a quorum of ‘Acceptors’. Acceptors will then decide if they want to make a ‘Promise’ which is a commitment to accepting that proposal in the future. If a quorum of promises is received by the Proposer, then it will assign a value to its proposal and will again send an ‘Accept Request’ out to a quorum of Acceptors. Finally, if enough ‘Accept’ messages are received then the Proposer can be certain that the value has been agreed upon by consensus.

This algorithm has been proved to be consistent but it also has a lot of complexity and a lot of different variants. The combination of different roles and states makes it easy to implement incorrectly. It is also important to consider that Paxos describes a ‘family’ of

algorithms, with some parts left deliberately unspecified, and choosing how to implement these is not a trivial decision. The final issue with Paxos is that it cannot guarantee progress. Whilst it enforces conditions which make it unlikely that progress will not be made, it is still theoretically possible for the mechanism to stall indefinitely.

Raft

Raft [10] is an algorithm that was designed to be equivalent and as efficient as Paxos, however, it also places a much greater emphasis on comprehensibility. It uses the notion of a *strong leader*, which is an elected server that has total control over which log entries are accepted. There are two other types of server, a *follower* and a *candidate*. Followers are completely passive, and only respond to requests from leaders and candidates. A candidate is a server that has put itself forward for election.

So how does the algorithm operate? Time in Raft is split into *terms* which are labelled by a monotonically increasing integer. Each term effectively signals a time period where a particular server is the leader. A term starts with a leadership election when a follower transitions into the candidate state, increments its term number and requests votes from other servers. It will wait for a majority of votes and then elect itself the leader, unless it times out or receives a message from another leader with a greater term number. Typically, these leader election processes are triggered when a follower does not receive a heartbeat message from the leader for longer than a given time. In the pathological case, the vote can be split between leaders, triggering a new election which is also split and so on. However, Raft uses randomised election timeouts to avoid this problem. Additionally, Raft also prescribes some restrictions on the servers which can be elected leader so as to avoid newly elected leaders overwriting previously appended log items.

Raft is a simple algorithm that is easy to understand, but it also guarantees the Log Matching Property that if two logs contain a log entry with the same index and term, then that entry, and all preceding entries will be identical.

3.3.2 A New Approach

Logan uses a mempool and a simple leader based algorithm to maintain consensus, and in this section, I will describe this approach. Because I have implemented a centralised algorithm, my specification has differed slightly from that presented in the Requirements Analysis to take into account differing functional requirements for a leader and for a participant. As I am assuming that all participating nodes can be trusted, it is possible to use a leader based approach without introducing security issues such as the ones tackled by the Proof of Work mechanism. This approach also reduces the potential complexity and overheads of implementing completely decentralised consensus. My approach makes use of the mergeable log data structure provided by Ezirmin, and the synchronisation module that allows updates from a log to be pulled into another. Finally, I have introduced the

notion of validation which allows both participants and leader nodes to accept or reject transactions depending on arbitrary conditions.

Leaders

Logan uses the notion of a *Strong Leader* similar to that used by the Raft protocol. The Leader is chosen statically in order to reduce the complexity of implementation, and it also has a statically defined list of `remotes` which specifies the location of all participants. The Leader is a node which will never actually request transactions to be added to the blockchain, its role is simply to periodically read requests from the mempools of Participants, validate them, and then add them to the blockchain. This blockchain can be read by any Participant, and is treated as the empirical source of which transactions have been committed and in which order. That is, any two nodes that read a copy of the blockchain, will always agree on content and ordering of log items in the blockchain, up until the end of the shortest copy (or both copies).

Listing 3.3: Leader Specification

```
module type I_LeaderConfig = sig
  type t
  module LogCoder: Participant.I_LogStringCoder with type t
    = t
  val remotes: string list
  val validator: (t list -> t list -> t list) option
end
module type I_Leader = sig
  val init_leader: unit -> (unit -> unit Lwt.t) Lwt.t
end
module MakeLeader (Config: I_LeaderConfig) : I_Leader =
  struct
    ...
  end
```

Listing 3.3 is the specification of the Leader module, which differs slightly from the specification presented in the Requirements Analysis. In particular, the `init_leader` function will perform an initialisation step, and then return a function which, when executed, will actually start the consensus process. Also of interest is the supplied `validator` function which will take a list of all committed transactions, a list of requested transactions and will return a filtered list of all valid transactions to be added.

Participants

Participants, in contrast to Leaders, can request transactions to be added to the blockchain. This is done by writing a transaction to a local mempool, which is then read by the leader. Validation can of this transaction can also happen at this stage, but it

cannot filter out all invalid transactions, and only exists to ease load on the leader.

Listing 3.4: Participant Specification

```

module type I_LogStringCoder = sig
  type t
  val encode_string: t -> string
  val decode_string: string -> t option
end
module type I_ParticipantConfig = sig
  type t
  module LogCoder: I_LogStringCoder with type t = t
  val leader_uri: string option
  val validator: (t list -> t -> bool) option
end
module type I_Participant = sig
  type t
  val add_transaction_to_mempool: t -> [> `
    Could_Not_Pull_From_Remote | `Validation_Failure | `Ok]
    Lwt.t
  val get_transactions_from_blockchain: int -> [> `Error | `
    Ok of t list] Lwt.t
  val get_all_transactions_from_blockchain: unit -> [> `
    Error | `Ok of t list] Lwt.t
end
module Make(Config: I_ParticipantConfig): I_Participant with
  type t = Config.t = struct
    ...
  end

```

Listing 3.4 is a specification that shines a light on the role of Participants and the functionality they provide. The module includes the ability to define custom data types that can be stored on the blockchain, to define how to validate transactions, to read from the blockchain, and to attempt to write to the blockchain by writing to a mempool.

Retrieving local updates

The mechanism used by the Leader to read mempool updates is different for the situations when the Participant is on the same machine, and when it is on a remote machine. Here, I will detail the process of reading updates from a participant on the same machine as the leader.

The leader will always maintain a cursor to the latest entry it has read from the participant mempool. When it attempts to find new updates, it will get a new cursor




Figure 3.4: Typical Participant mempool

to the latest element of the mempool. The leader can then compare these two cursors, and if the *latest known* cursor points to an entry with an earlier timestamp than the *latest* cursor, then it will add the *latest* item to an accumulator, and then repeat the process with a cursor pointing to the parent of the *latest* node. Finally, when the cursors match, the items in the accumulator are returned as new updates which can be added to the leader mempool and the *latest known* cursor is updated accordingly. It is important to note that no validation is performed at this stage as no item has been added to the blockchain yet. Figure 3.4 shows how unseen blocks are read sequentially from a mempool until the seen blocks are reached.

Retrieving remote updates

The previous section demonstrates how updates can be retrieved from a mempool on the same machine as the leader. In this section, I will highlight how updates are retrieved from mempools on remote machines.

Irmin provides a Sync module which allows for the histories of a local and a remote Store to be combined according to a custom merge function. Ezirmin logs build on this functionality by providing a merge strategy which uses the log entry timestamps to order log entries in a merge. Figure 3.5 demonstrates how this merge works in a situation where a number of blocks have been added to a single remote participant mempool. The first (and most naïve) approach I took to pulling updates from a number of participants, was to sequentially merge updates from all the participants into the leader's mempool. Participants also merge updates from the leader's mempool before adding to their own mempool to increase the shared history and decrease the work required for the leader to perform the merge. After pulling all updates, the leader will traverse its own mempool, finding new updates until it reaches the *latest known* cursor. Now, the updates can be validated and added to the blockchain, and the *latest known* cursor updated to the latest message in the mempool.

However, this approach is fundamentally flawed and causes many updates to be overlooked by the leader. This is because of the delays that occur when mempool updates are retrieved over the network from multiple workers. For example, Figure 3.6 demonstrates a situation where there are just two workers. In this case, updates from the first worker will be merged before the second, however, after the first merge has taken




Figure 3.5: Naïve approach to retrieving mempool updates

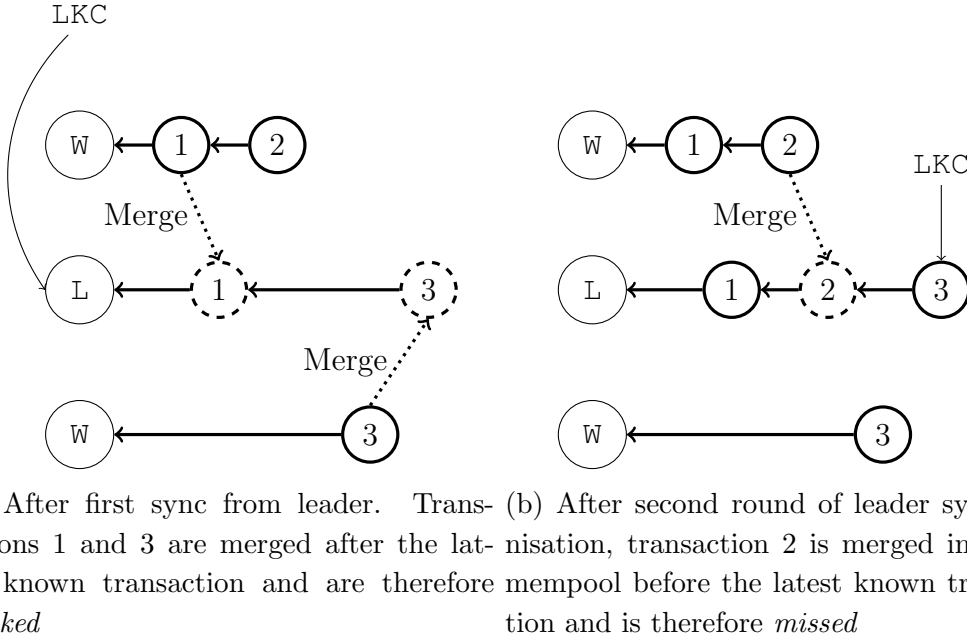


Figure 3.6: Missing mempool updates with a naïve consensus mechanism

place, the first worker may have added additional transactions. If these transactions are timestamped before the latest transaction in the second worker's mempool, then when the leader next polls for updates, transactions may be merged to a position in the mempool before the latest known cursor. This means that they will not be seen by the leader and therefore they will not be added to the blockchain. I shall refer to these transactions as *missed* transactions, and all others as *tracked* transactions.

In order to mitigate this problem, I first examined the nature of these *missed* transactions and noted the following properties:

1. Missed transactions must have been added during a merge. If they occurred before the merge had begun, then they would have been tracked. Alternatively, missed transactions must occur after the latest known cursor from the previous merge.
2. Missed transactions must have been added at a point in time earlier than the transaction pointed to by the latest known cursor. If they occurred at a later point in

time, then they would be tracked by future leader polls.

These lead very naturally onto a less naïve algorithm for retrieving updates which only adds updates that have existed for more than one ‘poll cycle’. Instead of maintaining a single cursor to the latest-known item, another cursor is now maintained to the previous latest-known item. I will refer to these as the `latest_known` and `previous_latest_known` items. After merging the newest set of updates, a leader can be sure that no more missed transactions will be added before the `latest_known` item, and after the `previous_latest_known` item.

Listing 3.5: Selecting new updates

```
let is_early_enough = is_earlier_or_equal scanning_cursor ~
  than:latest_known in
let is_late_enough = is_later scanning_cursor ~than:
  previous_latest_known in
match is_early_enough, is_late_enough with
| Some(false), Some(true) -> (*Item cannot be added to
  blockchain yet. Move scanning cursor back by one item
  and loop*)
| Some(true), Some(true) -> (*Item can be added to
  blockchain. Add item to accumulator, move cursor back
  by one item and loop*)
| _ ->(*Too far back in the mempool, so return item
  accumulator*)
```

Listing 3.5 shows a snippet of code from within a recursive function to get updates from the mempool, which can be added to the blockchain. This demonstrates a new approach which starts with a newly retrieved cursor, `scanning_cursor`, to the latest mempool element and then scans back through the mempool, adding valid items to an accumulator and eventually returning. On each loop, this algorithm will check if the item is valid, adding the item to the accumulator if this is the case, and then either return the accumulator or move the `scanning_cursor` to the next item back in the mempool. Every item from the `latest_known` item (inclusive) to the `previous_latest_known` item (exclusive) will be added to the blockchain. Any of the further updates are not safe to add, and will be postponed until the next poll. In the case that both cursors point to the same item (i.e. no new items were retrieved in the latest poll), no items are added as it is always assumed that the `latest_known` item is added in the previous poll. The cursors can now be updated as follows, where `get_new_latest_cursor` will get a cursor to the latest item in the leader’s mempool.

```
previous_latest_known_cursor := latest_known_cursor;;
latest_known_cursor := get_new_latest_cursor;;
```

This approach will not only pick up all the previously missed transactions, but all the tracked ones too. Figure 3.7 demonstrates how this approach would solve the problem

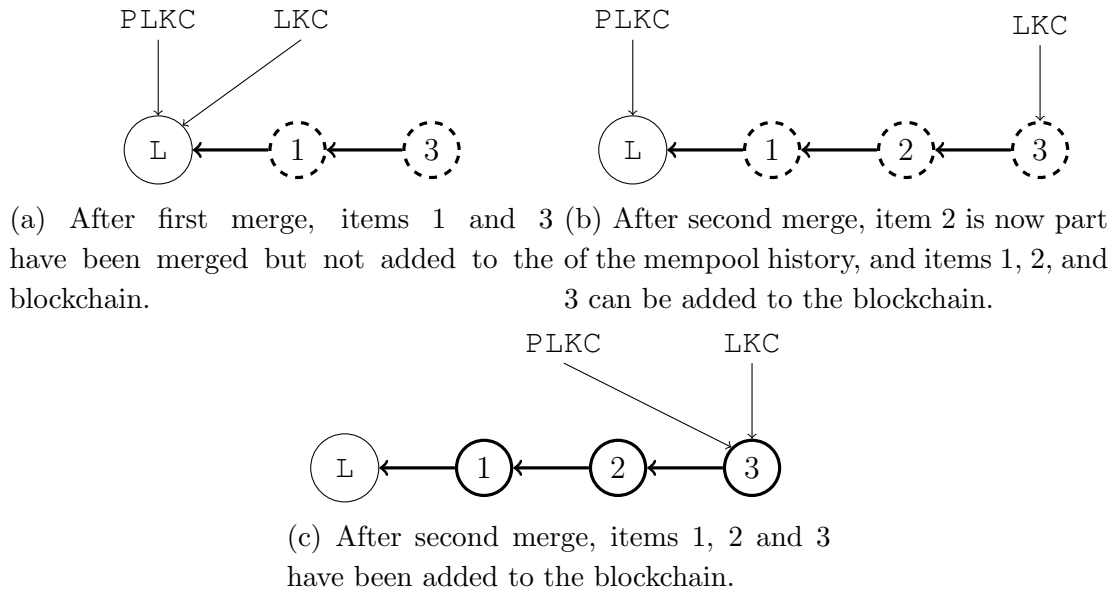


Figure 3.7: Advanced approach to retrieving mempool updates

presented in Figure 3.6.

At a first glance, it may seem sensible to use the `latest_known` cursor to get new updates from the mempool by just iterating back through the mempool from cursor. The problem with this approach is that the cursor is an abstraction of a tag to a block in the Irmin Block Store. This means that the history of the mempool according to that cursor is not changed by any subsequent merge operations. Consequently, any *missed* transactions will not be visible using this approach.

This now begs the question “How can the new approach, which uses out-of-date cursors, still be valid?”. The answer to this is that, the new approach only uses the cursors to perform timestamp comparisons. The result of these comparisons is the same whether a cursor to the items in the out-of-date mempool or a cursor to the items in the new mempool is used.

Now I have detailed all the work that I have completed to build the blockchain library. Transactions are sourced from the mempool and added when it can be certain that there will be no missed transactions. Transactions are added in order and never changed such that if a participant observes the state of the blockchain at any point, it will never observe a conflicting state in the future. The algorithm used by this project assumes that clock timestamps are synchronised across all devices. This may not be the case in an arbitrary situation, and consequently, in high load situations some transactions may not be recognised by the above algorithm. However, the system is not built for systems of such high load and I do not consider the missing of small numbers of transactions to be a large issue. In the worst case, an participant will be able to observe whether a transaction has not been added to the blockchain, and after a certain elapsed time period, can attempt again to add the transaction

Chapter 4

Evaluation

Metrics Used

In order to evaluate the performance of Logan, I have looked at two key metrics: latency and throughput. In a blockchain system, it is important that the latency is as low as possible. Latency dictates how long a user must wait before they can be sure that their transaction has been added. It is easy to see why this is important by thinking in terms of a cryptocurrency application. In this context long latencies will cause users to become irritated and potentially stop using your currency. Throughput is also a useful measure because it shows how well a system can cope with large loads. Again in the context of a cryptocurrency, it is important that lots of users should be able to make transactions at the same time. As with all distributed systems, these two metrics are inevitably linked, and higher system throughputs will lead to larger latencies. In the worst case, a large throughput will cause the system to become overloaded, latencies will rise and throughput will even decrease! I have looked at how the throughput applied to the system affects the latency of Logan transactions, in order to try to determine the optimal use case conditions for a Logan system.

I have also looked at how the latency of adding transactions varies with the size of the blockchain at a constant throughput. It is important that the system should be able to work well even when the size of the blockchain is large so that performance will not degrade over time.

4.1 Performance on a Local Machine

4.1.1 Single Participant

Figure 4.1 shows how the latency of transactions varies with throughput for a Participant node on the same machine as the Leader node. Latency stays at between 10-15ms up until the system reaches about 70 transactions per second. The set of data points which lie significantly above the straight line are points where the system has become overloaded. The rate of transactions has become so high that latencies have increased and the throughput has actually decreased again. This is a really interesting result because it




Figure 4.1: Latency variation with throughput on a single local Participant

shows that Logan’s logic functions well at throughputs of between 50 and 60 transactions per second, whilst maintaining low transaction latency. These figures are much better than my success criteria which suggested the system should be able to cope with around 2-3 transactions per second.

Figure 4.2 shows how the latency of Logan transactions varies as the size of the blockchain increases. The Figure shows clearly that the latency of adding a transaction is constant in the size of the blockchain. This is a very desirable trait as it means that Logan can be used for systems that run for indefinite periods of time, with arbitrary long blockchain histories. In the requirements analysis, I set out the criteria that Logan should be able to add transactions at constant latency and this result shows that I have been successful in achieving this criteria.

The graph also compares latencies when a simple, constant-time validator is used, which blindly accepts all transactions, and when no validation is performed. The result without validation was obtained by removing the validation code completely from the codebase and just accepting all transactions. Intuitively, these two results should be the same because they perform the same logical task, and this graph confirms this intuition. From this, we can conclude that if a user provides a constant time algorithm for validating transactions, then the overall time for adding transactions will remain constant.




Figure 4.2: Latency variation with blockchain size on a local Participant

4.1.2 Multiple Participants

4.2 Performance on Remote Machines

4.2.1 Single Remote Participant

4.2.2 Multiple Remote Participants




Figure 4.3: Logan performance with a single remote Participant

Chapter 5

Conclusion

Bibliography

- [1] Bitcoin energy consumption index. <https://digiconomist.net/bitcoin-energy-consumption>.
- [2] Ethereum white paper. <https://github.com/ethereum/wiki/wiki/White-Paper>.
- [3] Gnu make. <https://www.gnu.org/software/make/>.
- [4] jbuilder. <https://jbuilder.readthedocs.io/en/latest/overview.html>.
- [5] Lwt: A cooperative threads library for ocaml. <https://ocsigen.org/lwt/>.
- [6] GAZAGNAIRE, T. Introducing irmin: Git-like distributed, branchable storage. <https://mirage.io/blog/introducing-irmin>, 2014.
- [7] HABER, S., AND STORNETTA, W. S. How to time-stamp a digital document.
- [8] LAMPORT, L. The part-time parliament.
- [9] NAKAMOTO, S. Bitcoin: A peer-to-peer electronic cash system.
- [10] ONGARO, D., AND OUSTERHOUT, J. In search of an understandable consensus algorithm.
- [11] SIVARAMAKRISHNAN, K. Ezirmin: An easy interface to the irmin library. <http://kcsrk.info/ocaml/irmin/crdt/2017/02/15/an-easy-interface-to-irmin-library/>, 2017.
- [12] TORVALDS, L. Git: A distributed version control system, 2005.
- [13] YARON MINSKY, ANIL MADHAVAPEDDY, J. H. *Real World OCaml*. O'Reilly Media, Sebastopol, California, 2013.

Appendix A

Project Proposal

Building a Blockchain Library for OCaml

Charlie Crisp, Pembroke College

December 29, 2017

Project Supervisor: KC Sivaramakrishnan

Director of Studies: Anil Madhavapeddy

Project Overseers: Timothy Jones & Marcelo Fiore

Introduction

The blockchain, in its simplest form, is a tree-like data structure. Chunks of data are stored in 'blocks' which contain the hash of the contents of the previous block. This creates a 'blockchain' which can exhibit branching in the same way that a tree data structure can (see Figure 1). One of the most important features of a blockchain, is that a change in a block, will alter the block's hash, thereby altering all the future blocks in the chain. This makes it very easy to validate that the data in a blockchain is trustworthy, by verifying the hash in a block, is the same as the hash of it's parent's content.

Blockchain technology has generated a lot of interest in recent times, but




Figure 1: A typical blockchain structure [3]

mostly in the field of cryptocurrencies. With a simple Proof of Work consensus algorithm, the blockchain can be used to build a secure, distributed ledger of transactions. However, whilst the uses of the blockchain are far wider reaching than cryptocurrencies, progress outside of this field has been much slower.

I will build a pure OCaml, reusable blockchain library to allow the creation of distributed, secure ledgers, which are agreed upon by consensus. The library will allow users to create and add entries to a distributed blockchain ledger with just a few lines of code. The users will also be able to trust that entries in the blockchain are exactly replicated across all nodes in the network.

It will be built on top of Irmin [1] - a distributed database with git-like version control features. Being pure OCaml, the blockchain nodes can be compiled to unikernels or JavaScript to run in the browser. I will evaluate the blockchain by prototyping a decentralised lending library and evaluating the platform's speed and resilience.

Starting point

The project will build upon functionality provided by Irmin [1] which is a distributed database system. Irmin is fast, durable and has the branching capabilities which are required to build a blockchain.

Resources required

I will be using a Macbook provided by OCaml Labs [2] in order to develop the source code for the project. If the Macbook fails, then I will easily be able to transfer my work onto the MCS machines, as my project has no special requirements.

My work will also be backed up to a git repository hosted on GitHub and saved to a dedicated memory stick on a daily basis.

During the evaluation stage I will be running my platform on different cloud based devices and/or Raspberry Pi's. There are many possible providers for cloud computing, including Amazon Web Services and Microsoft Azure. OCaml Labs [2] will provide the necessary funds to acquire these resources.

Background

Consensus

Consensus is a group process where a network of nodes will reach a general agreement. There are different ways of achieving consensus but here are some of the most common:

1. **Proof of Work:** Trust is given to nodes which can prove that they have put in computational work. This is the consensus mechanism used by Bitcoin.
2. **Proof of Stake:** Nodes are selected to validate blocks based on their stake in the blockchain. There are few variations on this algorithm which introduce notions such as delegation or anonymity.
3. **Raft Consensus:** A leader is elected and acts as a governing authority until it fails or disconnects, whereupon a new leader is elected.

Work to be completed

The work for this project will be split into the following major parts.

1. Design and build a module to allow nodes to create and maintain a blockchain ledger. This will include allowing nodes to add blocks to the chain and to form new branches.

2. Design and build a module to allow nodes to interact over a network and to achieve consensus. As highlighted the Background section, there are many different ways to achieve consensus, and a large part of this work will be to determine which method is most suitable. This decision will take into account a method's failure tolerance in terms of nodes failing and network failure, as well as general speed and any requirements (e.g. computational work for a Proof of Work algorithm).
3. Design an application using these modules. This will take the form of a book lending platform where nodes will be able to register books and lend them to other nodes in the network. This application has been chosen, because the blockchain library should allow for typically centralised applications to be created in a decentralised way. It will also allow for testing of critical features, for example, books should never be 'doubly-spent', i.e. if one user believes they have ownership of a book, then no other user will think the same.
4. Design an evaluation program to simulate different load on the lending platform. This will be run in different configurations in order to measure the performance of the platform.

Evaluation metrics and success criteria

I will consider the project to be a success if the following criteria are achieved:

1. Nodes in the network are able to connect and communicate information.
2. Nodes are able to achieve consensus about the state of the distributed ledger.
3. Nodes are able to reconnect after being individually disconnected.
4. Nodes are able to re-converge after a network partition.

In order to evaluate the performance of the system, I will measure the *throughput* and *speed* of transactions of the book lending platform. Throughput will be measured in transactions per second, and speed will be quantified as time taken to complete a transaction. I will evaluate how these properties vary with respect to the following metrics:

1. **Number of nodes:** I will scale the number of nodes in the network between the range of 2 and 5.
2. **Rate of transactions:** I will vary the number of transactions made per second.

Should I achieve and be able to measure the above criteria within the time frame of my project, I will further test system against the following metrics:

1. **Network latency between nodes**
2. **Network bandwidth of nodes**

Timetable

1. **Michaelmas Weeks 2-4** (12/10/17 - 01/11/17):
Set up an environment for developing OCaml and familiarise myself with the language and it's module system. This is important because the blockchain library needs to be reusable, and therefore well isolated.
2. **Michaelmas Weeks 5-6** (02/11/17 - 15/11/17):
Familiarise myself with Irmin and it's data structures. This is important as I have never used the library before, but it will be used to build the blocks in the blockchain library. In this time I will also begin to design the API of my library.
3. **Michaelmas Weeks 7-8** (16/11/17 - 29/11/17):
Finalise the API and start to build the module for creating and interacting with a distributed ledger. This will also involve investigating which hashing algorithms can be used to form the blockchain data structure.
4. **Christmas Vacation** (30/11/17 - 17/01/18):
Finalise the API of the module for achieving consensus between multiple nodes. This work will also include investigating different methods of consensus and their suitability for my project.
5. **Lent Weeks 1-2** (18/01/17 - 31/01/18):
Build the module for achieving consensus between modules. I will also start work on an lending library application which will be used to evaluate the performance of the blockchain library.

6. **Lent Weeks 3-4** (01/02/18 - 14/02/18):
Finish work on the lending library application and install it on a number of Raspberry Pi and/or cloud based devices. I will also begin work on my dissertation and I aim to complete the Introduction and Preparation chapters.
7. **Lent Weeks 5-6** (05/02/18 - 28/02/18):
Evaluate the performance of the platform by simulating load from each of the devices and measuring the speed of transactions. A stretch goal for this period is also to evaluate a range of further metrics. Additionally I will continue work on my dissertation and aim to complete the Implementation chapter.
8. **Lent Weeks 7-8** (01/03/18 - 14/03/18):
Finish a first draft of my the dissertation by writing the Evaluation and Conclusion chapters. I will also send the dissertation to reviewers to get feedback.
9. **Easter Vacation** (15/03/18 - 25/04/18):
With a first draft of the dissertation completed, I will use this time to review the draft and to make improvements. I will also incorporate feedback from reviewers, and complete the Bibliography and Appendices chapters.
10. **Easter Weeks 1-2** (26/04/18 - 09/05/18):
Conclude work on dissertation by incorporating final feedback from reviewers.
11. **Easter Week 3-Submission Deadline** (10/05/18 - 08/05/18):
I aim to have completed the dissertation by this point, and to be focusing on my studies. However, this time may be needed to make any final changes.

References

- [1] Irmin - A pure OCaml, distributed database that follows the same design principles as Git.
<https://github.com/mirage/irmin>

- [2] OCaml Labs - An initiative based in the Computer Laboratory to promote research, growth and collaboration within the wider OCaml community
<http://ocamlabs.io/>
- [3] Image of blockchain data structure from Wiki Commons.
<https://commons.wikimedia.org/wiki/File:Blockchain.png>